

## Check list

- [ ] Microphone turned on
- [ ] Zoom room open
- [ ] MyBinder launched
- [ ] Sound Volume on

# SSIE 419/519: Applied Soft Computing

Lecture 01: Introduction & Word Embeddings

Sadamori Kojaku



# Binghamton University

EngiNet™

# State University of New York



## **WARNING**

**All rights reserved. No part of the course materials used in the instruction of this course may be reproduced in any form or by any electronic or mechanical means, including the use of information storage and retrieval systems, without written approval from the copyright owner.**

**©2025 Binghamton University  
State University of New York**



## Contact Information:

EngiNet Office Staff:

Janice Kinzer

Email: [enginet@binghamton.edu](mailto:enginet@binghamton.edu)

Phone: 1-800-478-0718 or 607-777-4965

Media Production Operator: Ryan Luo

Instructor: Sadamori Kojaku

Email: [skojaku@binghamton.edu](mailto:skojaku@binghamton.edu)

# Course Overview

- **Instructor:** Sadamori Kojaku (幸若完壮)
- **Email:** [skojaku@binghamton.edu](mailto:skojaku@binghamton.edu)
- **Office Hours:** Tue & Thu 15:00-17:00
- **Lecture Note:** <https://skojaku.github.io/applied-soft-comp>
- **GitHub:** <https://github.com/skojaku/applied-soft-comp>

# Why Soft Computing?





103 Emerging  
Jason Allen  
Pueblo West  
Théâtre D'opéra Spatial

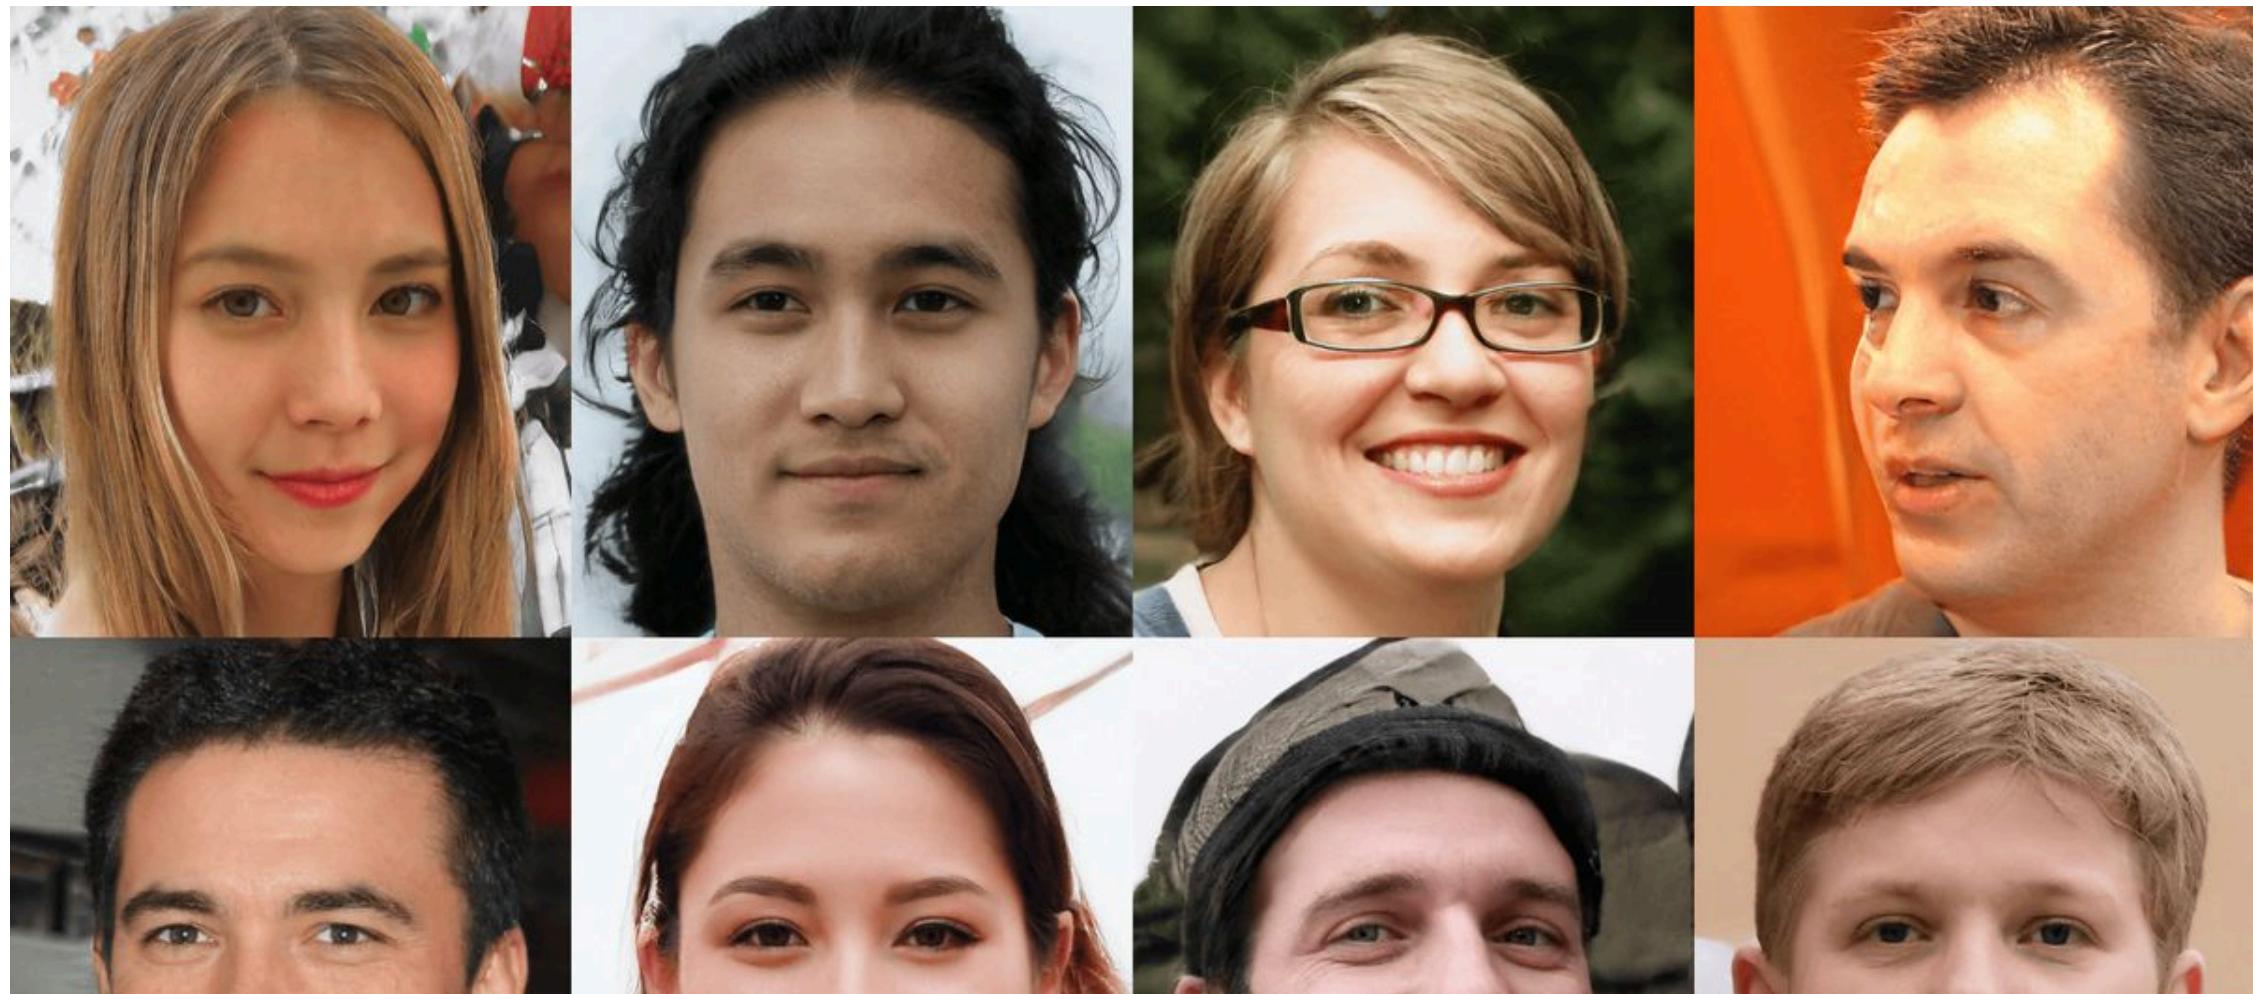
\$750  Colorado State Fair



bioRxiv preprint doi: <https://doi.org/10.1101/212022>; this version posted March 2, 2014. The copyright holder for this preprint (which was not certified by peer review) is the author/funder, who has granted bioRxiv a license to display the preprint in perpetuity. It is made available under a [aCC-BY-ND 4.0 International license](https://creativecommons.org/licenses/by-nd/4.0/).

# People that do not exist

[ThisPersonDoesNotExist.com](https://ThisPersonDoesNotExist.com)



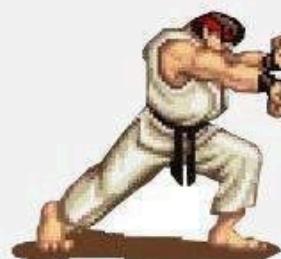




## How is it possible?

# 60s-80s: Expert System

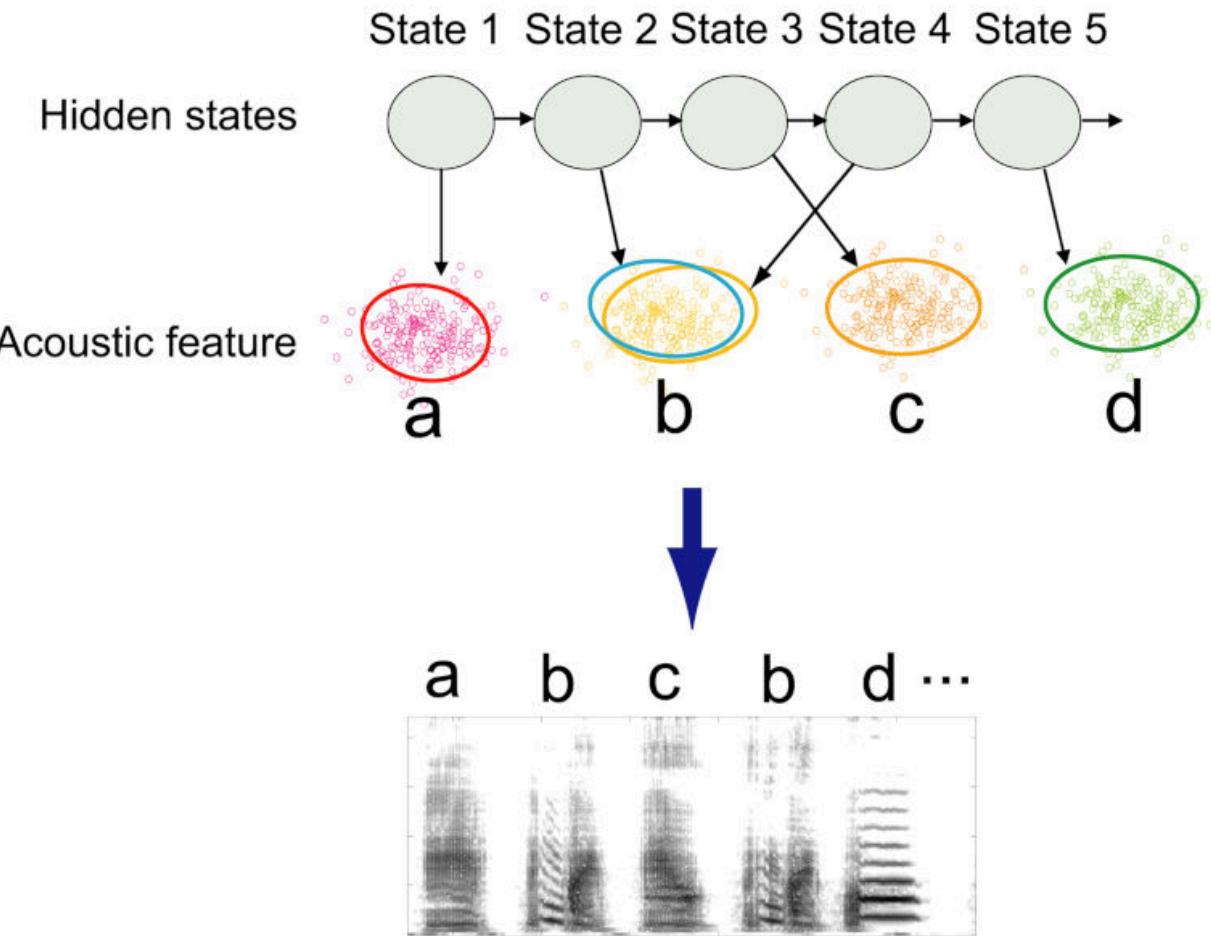
- Explicit rules
- Examples:
  - MYCIN (1976) for medical diagnosis
- Issues:
  - No adaptability
  - Need a lot of rules for complex tasks



```
function register()
{
    if (!empty($_POST)) {
        $msg = '';
        if ($_POST['user_name']) {
            if ($_POST['user_password_new']) {
                if ($_POST['user_password_new'] === $_POST['user_password_repeat']) {
                    if (strlen($_POST['user_password_new']) > 5) {
                        if (strlen($_POST['user_name']) < 65 && strlen($_POST['user_name']) > 1) {
                            if (preg_match('/^a-zA-Z\d{2,64}$/', $_POST['user_name'])) {
                                $user = read_user($_POST['user_name']);
                                if (!isset($user['user_name'])) {
                                    if ($_POST['user_email']) {
                                        if (strlen($_POST['user_email']) < 65) {
                                            if (filter_var($_POST['user_email'], FILTER_VALIDATE_EMAIL)) {
                                                create_user();
                                                $_SESSION['msg'] = 'You are now registered so please login';
                                                header('Location: ' . $_SERVER['PHP_SELF']);
                                                exit();
                                            } else $msg = 'You must provide a valid email address';
                                        } else $msg = 'Email must be less than 64 characters';
                                    } else $msg = 'Email cannot be empty';
                                } else $msg = 'Username already exists';
                            } else $msg = 'Username must be only a-z, A-Z, 0-9';
                        } else $msg = 'Username must be between 2 and 64 characters';
                    } else $msg = 'Password must be at least 6 characters';
                } else $msg = 'Passwords do not match';
            } else $msg = 'Empty Password';
        } else $msg = 'Empty Username';
        $_SESSION['msg'] = $msg;
    }
    return register_form();
}
```

# 90s-2000s: Statistical Learning

- Allow some randomness in real-world data
- Examples:
  - Support Vector Machine
  - Random Forest
- Issues:
  - Need a lot of data
  - No cross-domain generalization



# 2010s-2020s: Deep Learning



- Learn patterns from examples
- Adapt to new situations
- No need for explicit rules
- Handle complexity naturally

# Neural networks are not new

- Warren McCulloch and Walter Pitts to develop the concept of the McCulloch-Pitts (MCP) neuron in 1943.
- Frank Rosenblatt later introduced the perceptron learning rule for the MCP.



Perceptron Research from the 50's & 60's, clip



## Neural networks were virtually dead until 2012

- Computing limitations - computers were too slow to train meaningful neural networks
- Data scarcity - Neural networks require massive datasets that weren't available
- Vanishing gradient - Prevented training deep architectures effectively
- Better alternatives (e.g., SVMs) outperformed neural networks with less overhead
- AlexNet (2012) represents a breakthrough in deep learning

# Language Understanding: ChatGPT

- Natural conversation abilities
- Multiple capabilities:
  - Poetry and creative writing
  - Technical explanations
  - Code assistance
  - Understanding context and humor



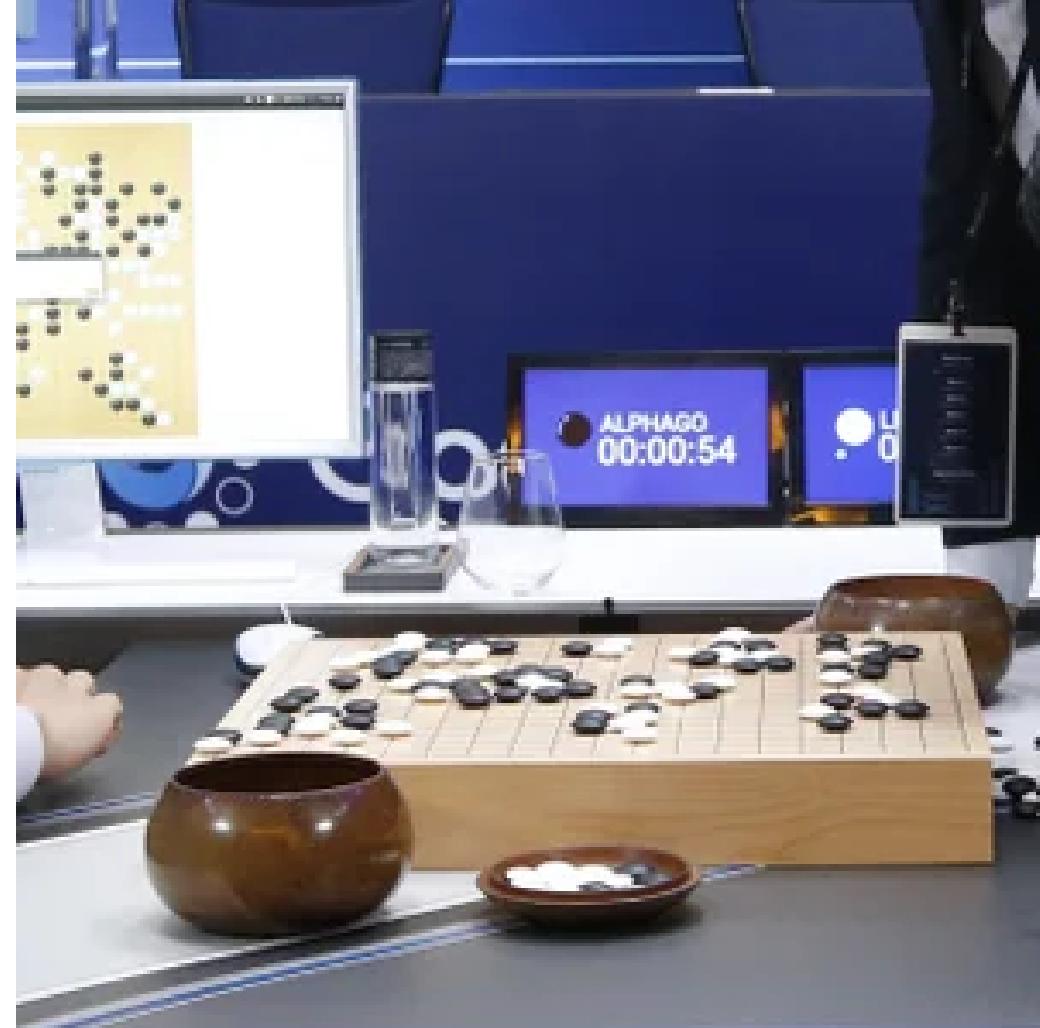
# Scientific Breakthrough: AlphaFold

- Solved 50-year protein folding problem
- Near-perfect accuracy
- Accomplished in days what took months in labs
- Revolutionary impact on biology and medicine



# Game AI: The Divine Move 🎮

- AlphaGo vs Lee Sedol
- Famous "Move 37"
- Demonstrated creative thinking
- Showed AI can surpass human intuition



AlphaGo



Lee Sedol



# Video Generation: Sora 🎬

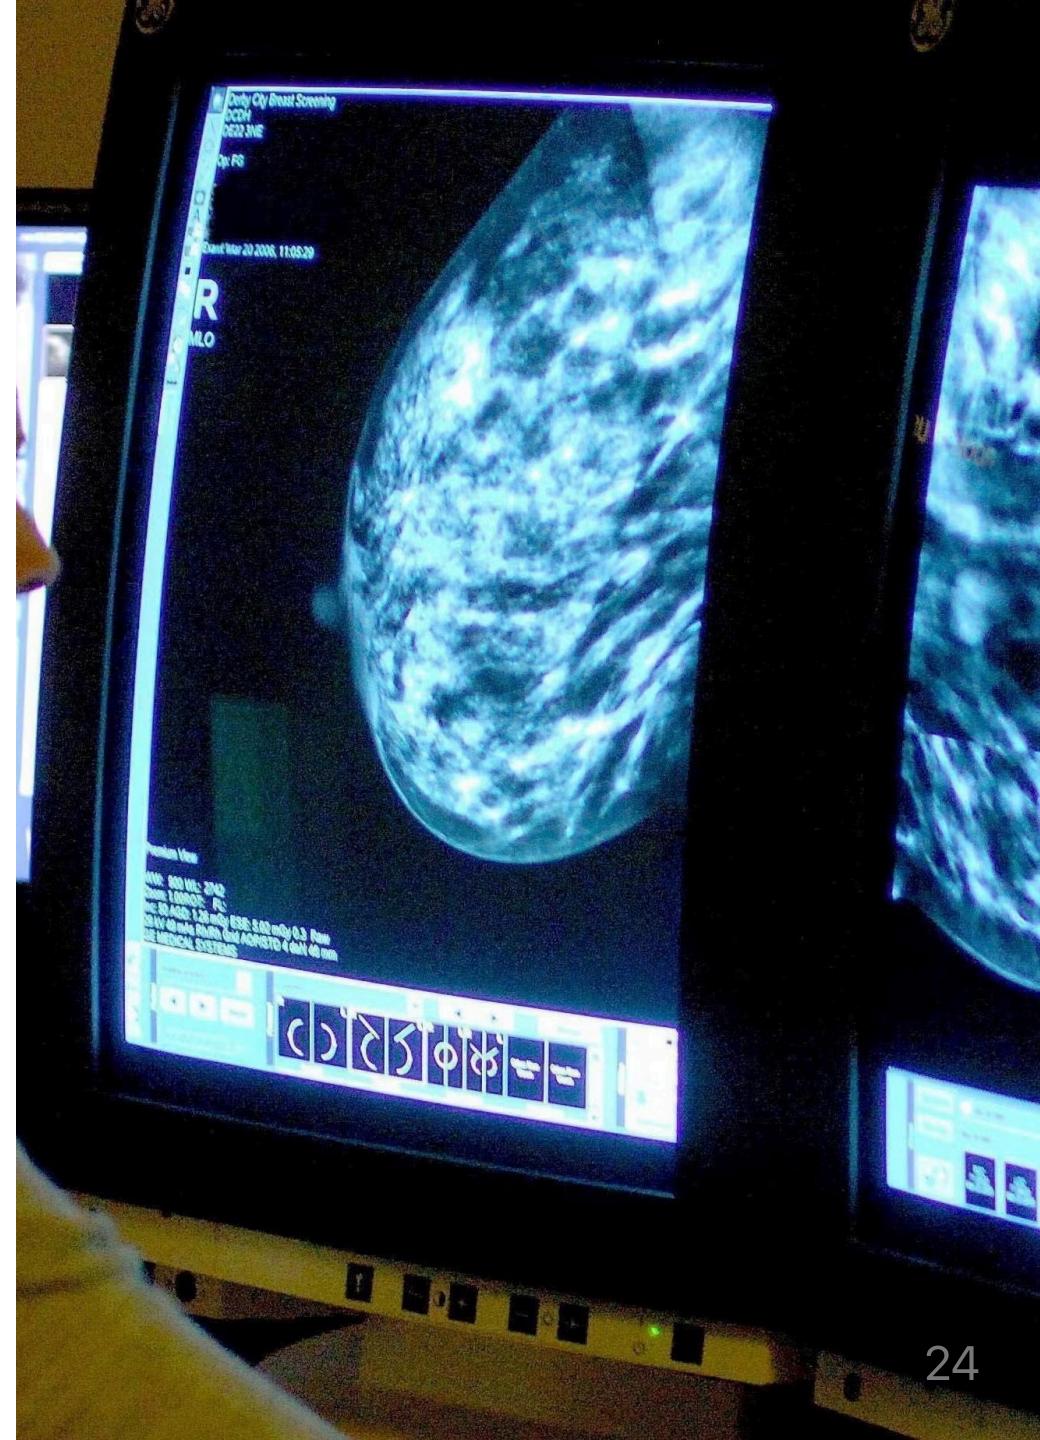
- Text-to-video technology
- 60-second realistic videos
- Physics-accurate scenes
- Multiple moving elements
- Photorealistic quality



# Medical Diagnosis: AI Assistant



- Surpasses human accuracy in cancer detection
- Reduces:
  - Missed cases
  - False alarms
- Augments doctor's capabilities



# Autonomous Systems: Self-Driving



- Real-time sensor processing
- Faster than human reactions
- Constant alertness
- Improved safety in many conditions



# Why This Matters

- Unprecedented pace of breakthroughs
- Solving decades-old problems
- Surpassing human capabilities
- Transforming multiple fields:
  - Healthcare
  - Transportation
  - Scientific research
  - Creative arts

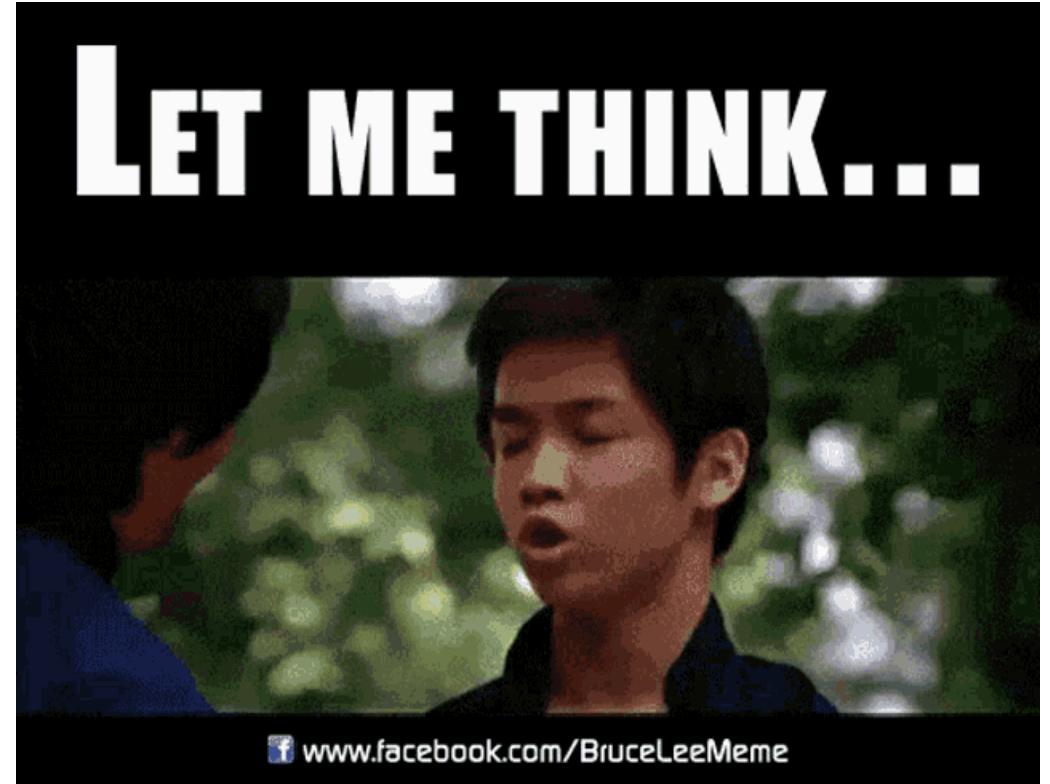
**Join the revolution in applied soft computing!**

# About this course

# Course Structure

"Don't think! Feeeeeeel" - Bruce Lee

-  Lectures
-  Hands-on exercises
-  Weekly quizzes
-  Biweekly coding assignments
-  Final project
-  Exam



# Final Project

- Individual project (30% of grade) 
- Timeline 
  - 03/09: Project Proposal
  - 05/06: Project Presentation
  - 05/09: Project Final Paper
- Requirements 
  - Apply concepts to real problem 
  - Show course integration 
  - Clear presentation 

# Exam

-  Final exam on all topics (weight: 30%)
-  During exam week
-  Theory + practical problems
-  Apply concepts to real scenarios
-  Review sessions before exam

# Weekly Quiz on Brightspace

-  Quizzes: A tool to identify misconceptions (weight: 20%)
-  Covers previous week's topics
-  Deadline: before final exam
-  Unlimited attempts until correct

# Assignment

-  One assignment per module (weight: 20%)
-  Coding exercises
-  Autograded assignments
-  Deadline: before final exam
-  Unlimited attempts until correct

# Lecture note

-  [Interactive Jupyter book](#)
-  Run code directly on the page
  -  First-time loading may take 2-3 mins
-  Or download as Jupyter notebook
  -  Use on cloud (Google Colab, Kaggle) or locally
  -  Install packages from `environment.yml` for local use
  - See [The course GitHub repo](#) for details
- Slides: [The course GitHub repo](#)

# Policy

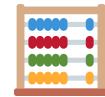
-  3-credit course: 6.5+ hours of work/week outside class
-  AI tools allowed for learning, but cite if used in assignments
-  Back up all data and code (loss not an excuse for late work)
-  Accommodations available for students with disabilities
-  Zero tolerance for academic dishonesty

# Questions?

# Teaching computers how to understand words



# The Challenge: Teaching Computers Language



- Computers only understand numbers
- Words need to be translated into numerical form
- Early approach: One-hot encoding
  - Each word gets a unique binary vector
  - Example: cat  $\rightarrow$  [1,0,0], dog  $\rightarrow$  [0,1,0]
- Problem: No semantic meaning captured

## One-hot encoding

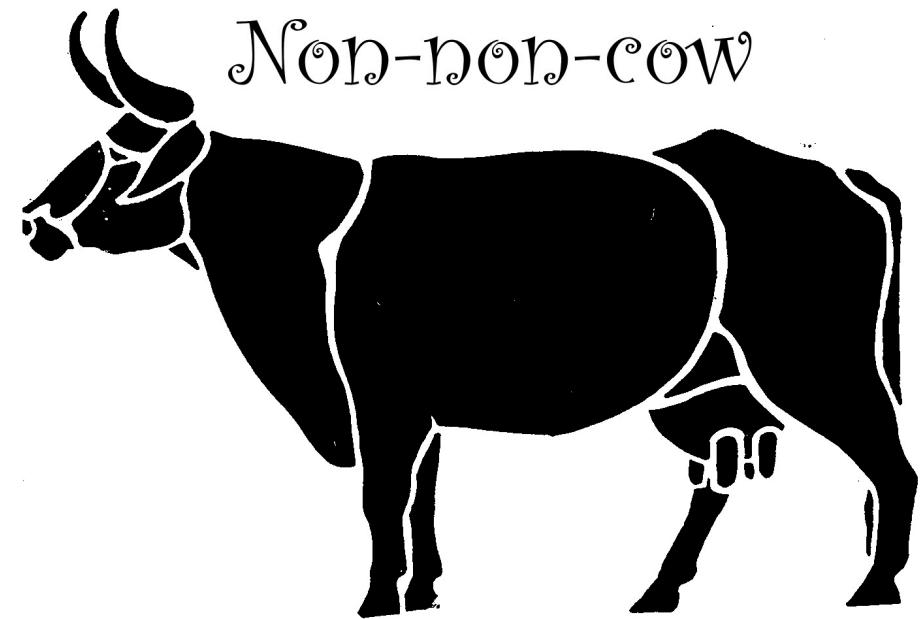
	cat	mat	on	sat	the	
the	=>	0	0	0	0	1
cat	=>	1	0	0	0	0
sat	=>	0	0	0	1	0
		...		...		

# Distributional Hypothesis

Words that appear in similar contexts  
have similar meanings

# A trivia

- Ancient Buddhist concept of Apoha (5th-6th century CE):
  - We understand concepts by what they are not
  - Example: A "cow" is defined by everything that is not a cow
- Parallels with modern distributional semantics:
  - Both define meaning through relationships between concepts
  - Words understood by their contrasts with other words



## Another trivia

My grandma called me in many different ways

- Sadamori
- My son
- My nephew
- My niece
- My cousin
- Her dog

# How can we build such a distributional representation?

Pen and Paper exercise

# From Words to Numbers: TF-IDF



- Term Frequency (TF):

$$TF(t, d) = \frac{\text{count of term } t \text{ in doc } d}{\text{total terms in doc } d}$$

- Inverse Document Frequency (IDF):

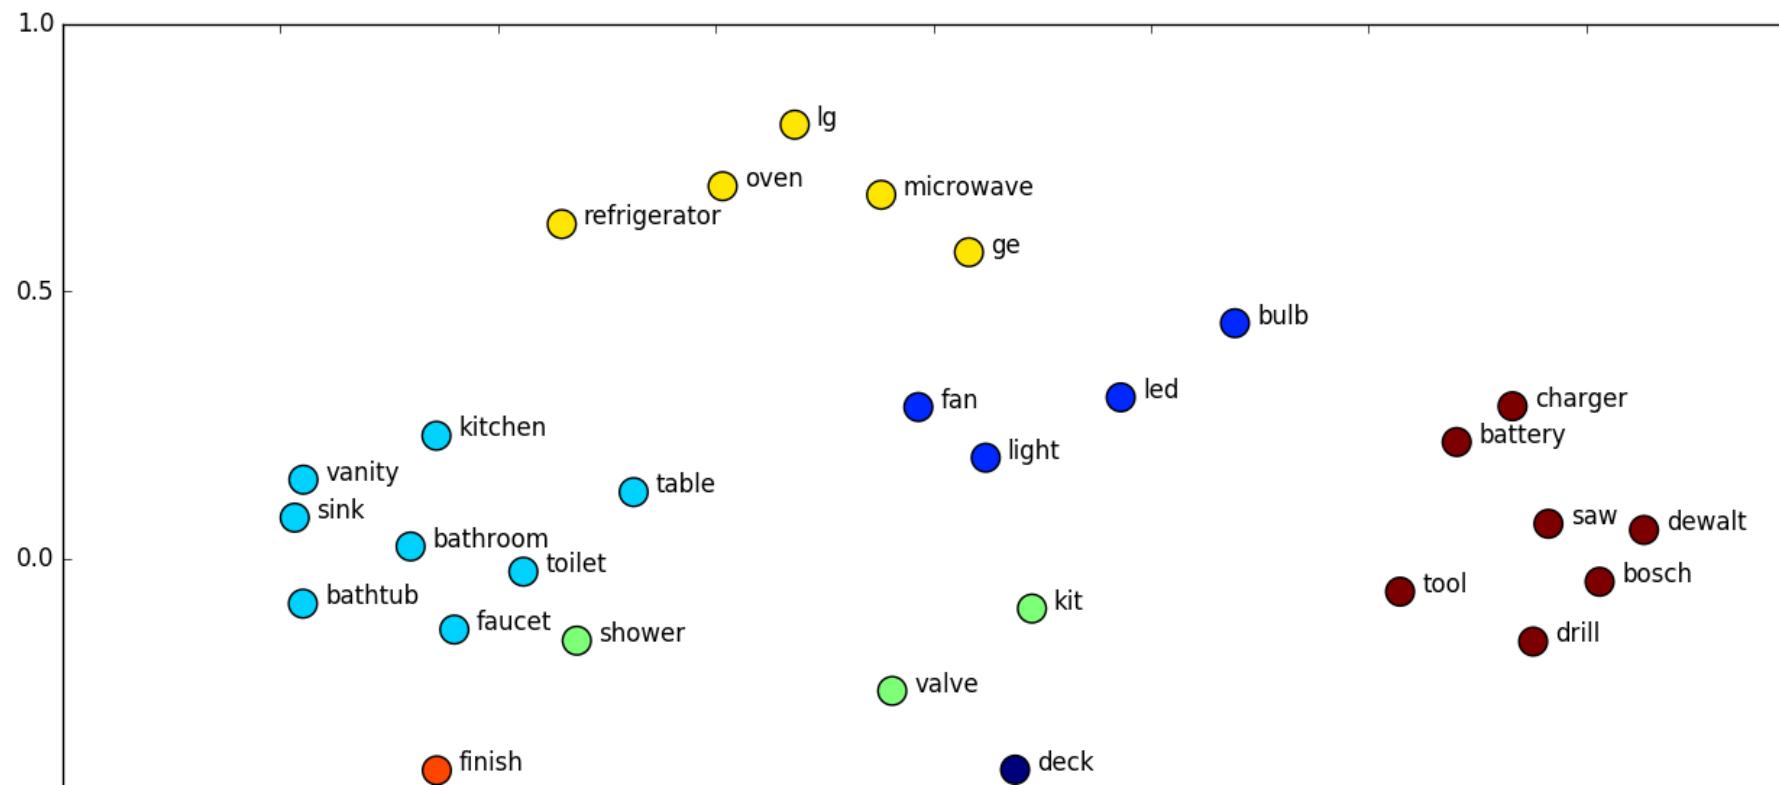
$$IDF(t) = \log \frac{\text{total documents}}{\text{docs containing } t}$$

- Combined score:

$$TF-IDF(t, d) = TF(t, d) \times IDF(t)$$

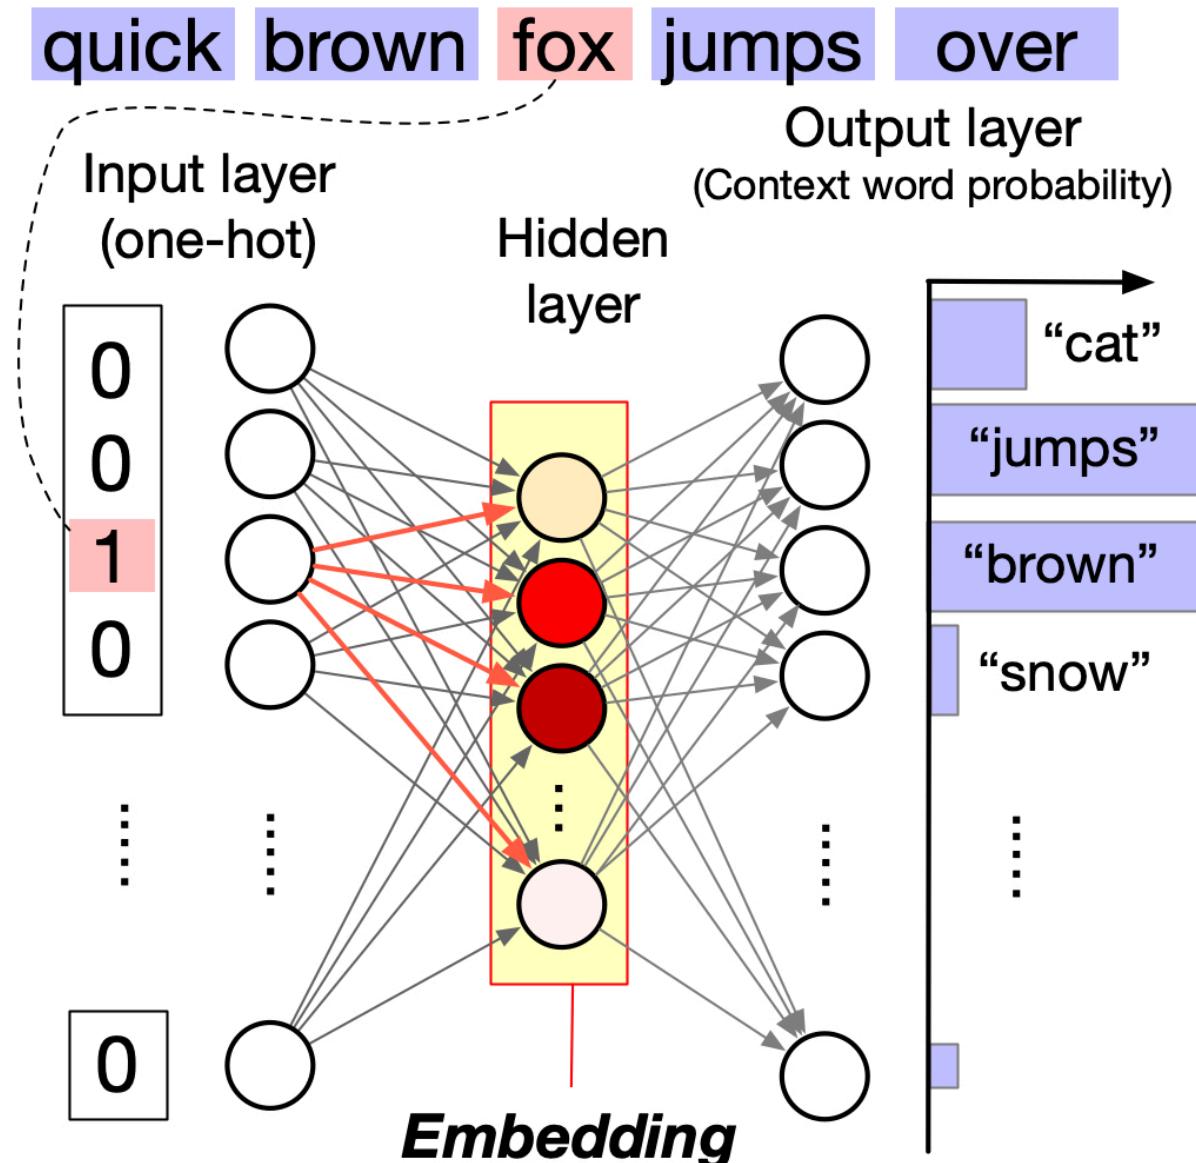
# The Distributional Hypothesis

"You shall know a word by the company it keeps"  
~Words appearing in similar contexts have similar meanings~



# Word2Vec: Neural Word Embeddings

1. Consist of one hidden layer without non-linear activation
2. Output layer is a softmax layer
3. Two models: Skip-gram and CBOW depending on the input and output



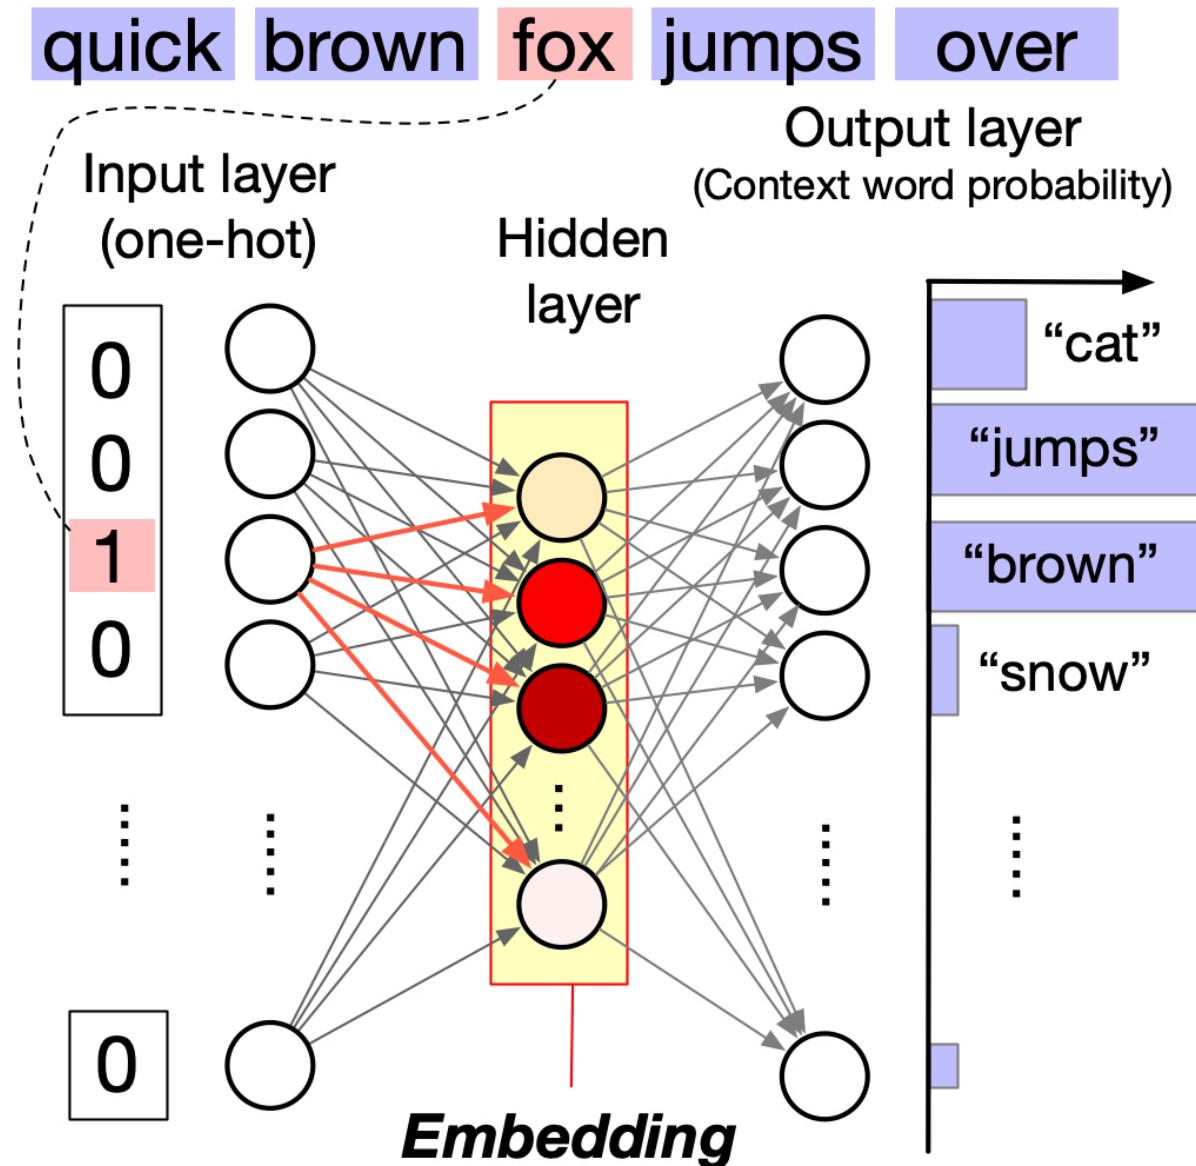
## Model (Skip-gram)

$P(\text{context word } j \mid \text{center word } i)$

$$= \frac{1}{Z} \exp(v_j^T u_i)$$

Each word has **two vectors**,  $v$  and  $u$ :

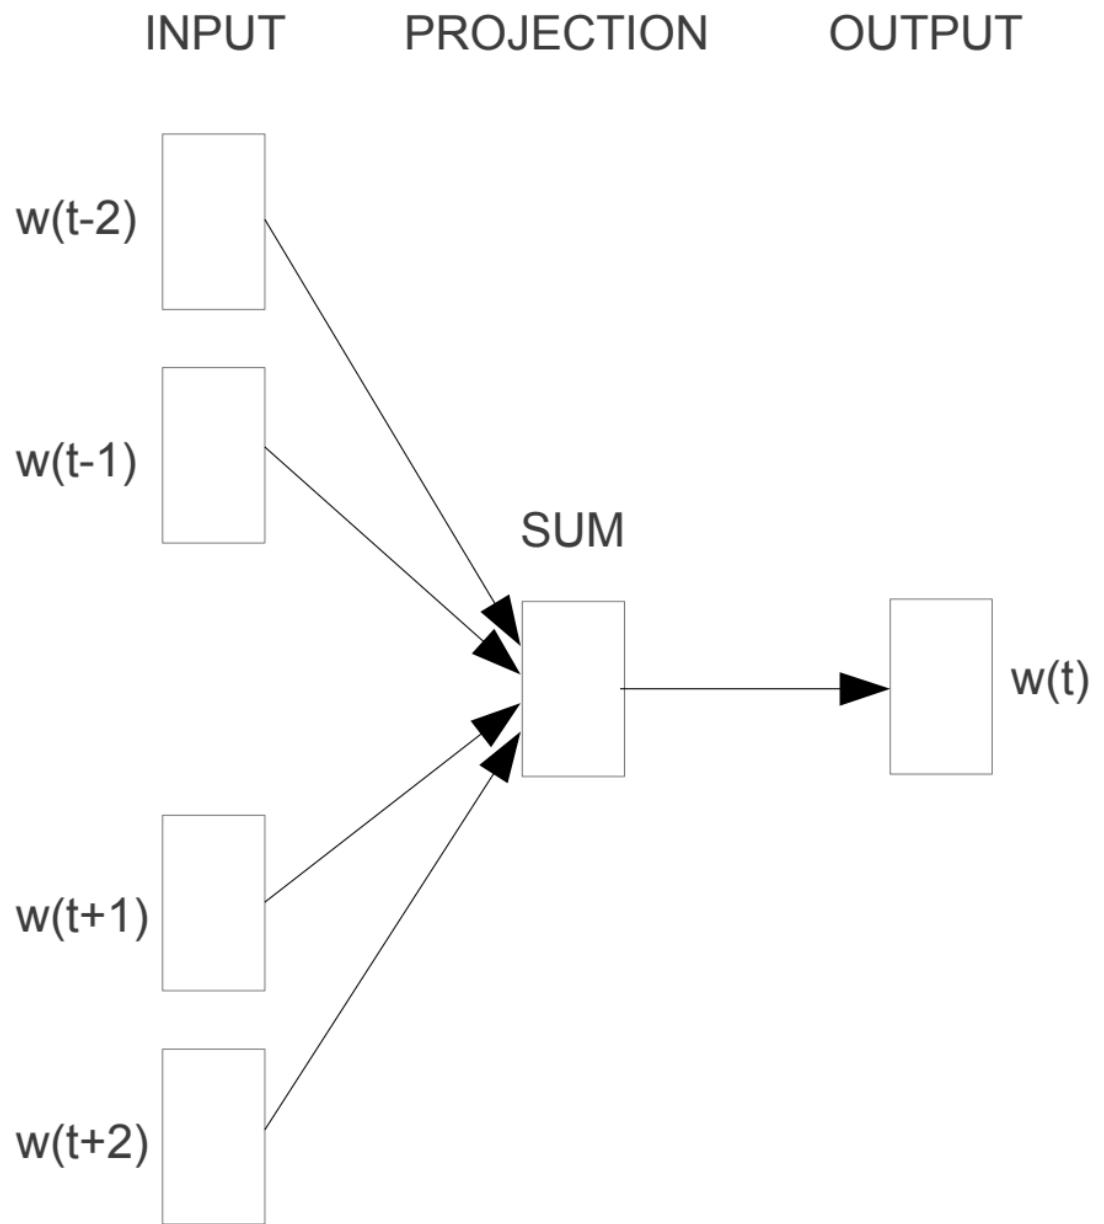
- In-vector:  $u_i$  represents word  $i$  as a center word
- Out-vector:  $v_j$  represents word  $j$  as a context word
- $Z$  is normalization constant



## Model (CBOW)

$$P(w_c|w_1, \dots, w_C) = \frac{\exp(v_{w_c}^T \bar{v})}{\sum_{w \in V} \exp(v_w^T \bar{v})},$$

where  $\bar{v} = \frac{1}{C} \sum_{i=1}^C v_{w_i}$  is the average of the context word vectors.



# Matrix Factorization Connection 12 34

Word2vec implicitly factorizes a Pointwise Mutual Information (PMI) matrix:

$$M_{ij} = \log \frac{P(w_i, w_j)}{P(w_i)P(w_j)}$$

Properties:

- Low when words appear independently
- High when words co-occur frequently
- Similar effect to TF-IDF normalization

Word embeddings preserve PMI values:

$$v_{w_i}^\top v_{w_j} \approx M_{ij}$$

Key insight: Words that frequently appear in similar contexts will have similar embeddings!

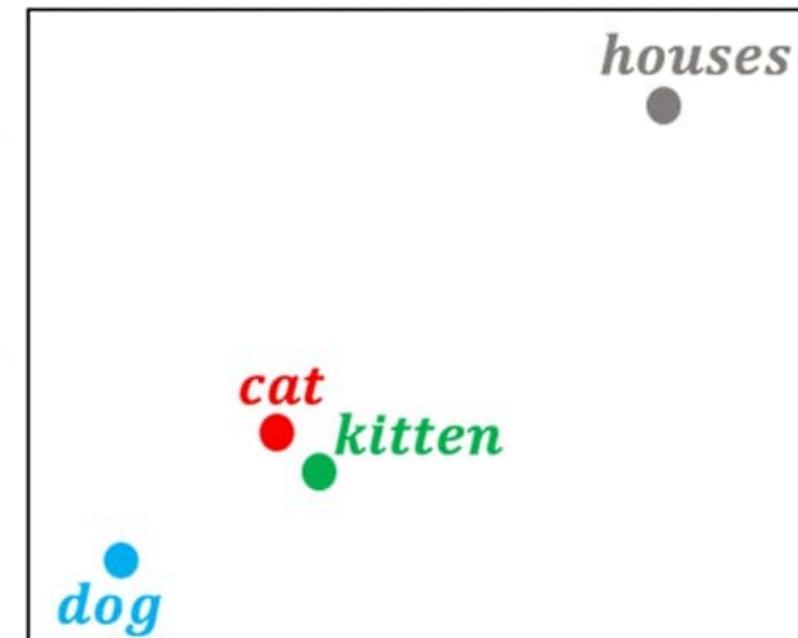
# GloVe: Global Vectors

GloVe explicitly factorizes PMI matrix:

$$M_{ij} = \log \frac{P(w_i, w_j)}{P(w_i)P(w_j)}$$

	living	being	feline	human	gender	royalty	verb	plural
<b>cat</b> →	0.6	0.9	0.1	0.4	-0.7	-0.3	-0.2	
<b>kitten</b> →	0.5	0.8	-0.1	0.2	-0.6	-0.5	-0.1	
<b>dog</b> →	0.7	-0.1	0.4	0.3	-0.4	-0.1	-0.3	
<b>houses</b> →	-0.8	-0.4	-0.5	0.1	-0.9	0.3	0.8	

Dimensionality  
reduction of  
word  
embeddings  
from 7D to 2D



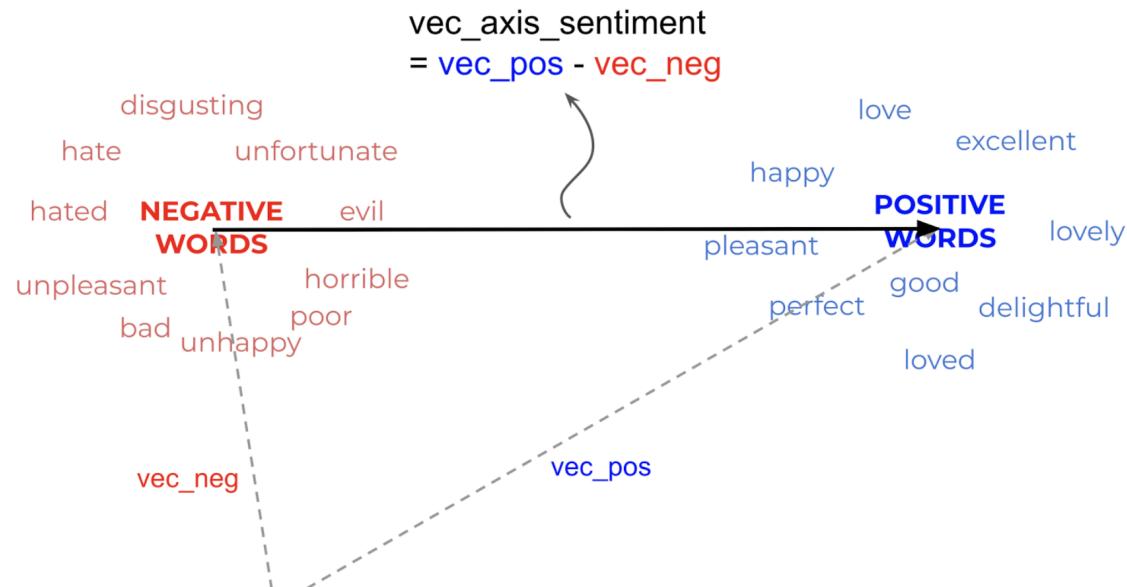
**Let's build the first word embedding using TF-IDF and word2vec**

[Jupyter notebook](#)

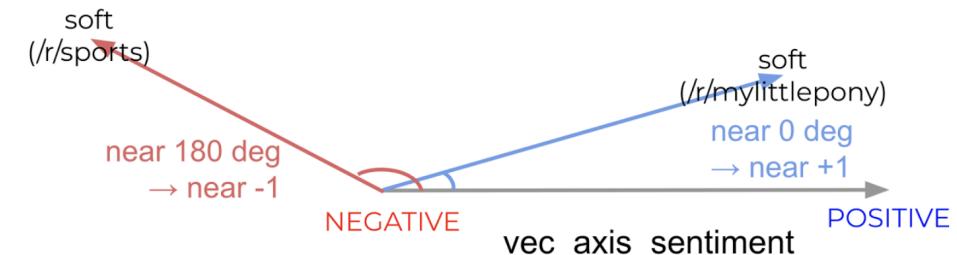
# SemAxis: Understanding Word Relationships

- Identify two semantically interpretable clusters
- SemAxis: An axis between the cluster centroids, e.g., good-bad, soft-hard
- Project word vectors onto the SemAxis

Step 1: Define a semantic axis



Step 2. Locate words on the semantic axis



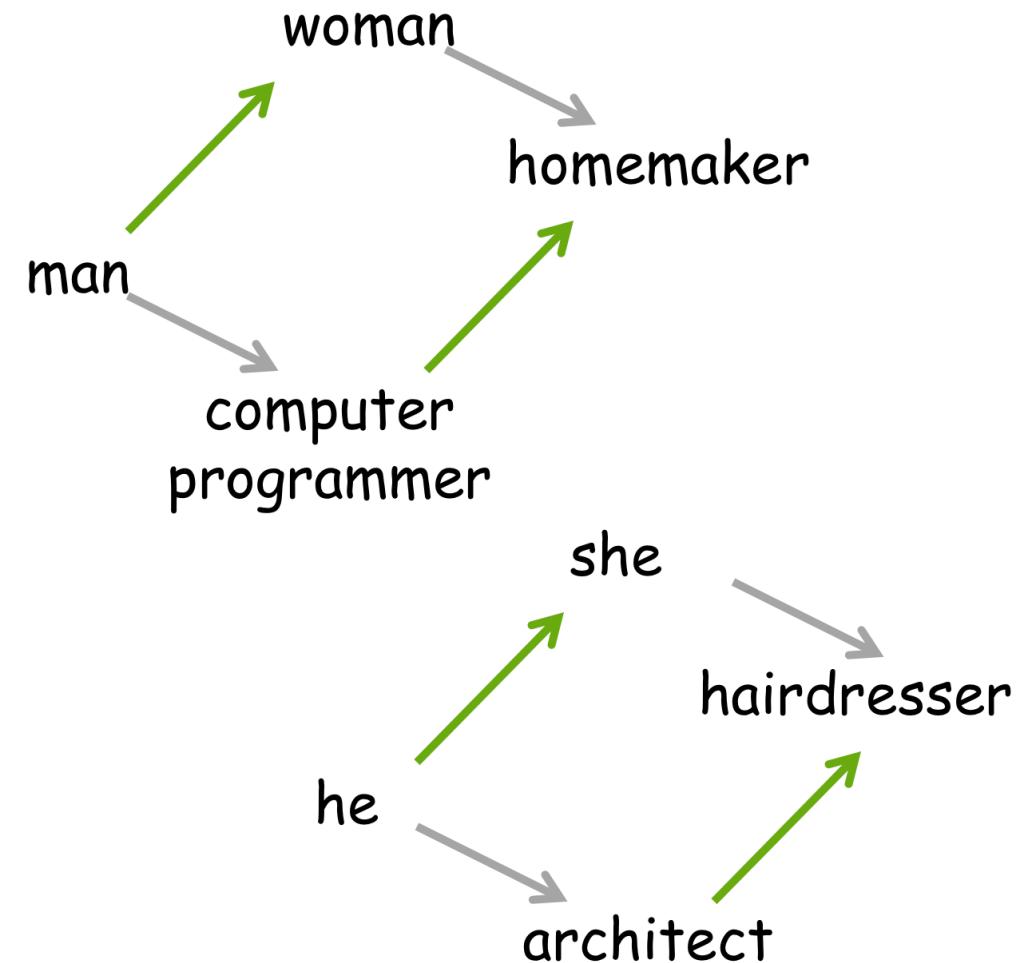
Compute a cosine similarity between a word vector and axis vector

# Bias in Word Embeddings



Gender bias example:

- man : doctor :: woman : nurse
- he : programmer :: she :  
homemaker



# Let's build SemAxis

[Jupyter notebook](#)

# Doc2Vec: From Words to Documents



Two models:

1. PV-DM (Distributed Memory)
2. PV-DBOW (Distributed Bag of Words)

## Doc2Vec (PV-DM)

PV-DM probability:

$$P(w_i | w_{i-k}, \dots, w_{i-1}, d) = \frac{\exp(u_{w_i}^T h)}{\sum_{w \in V} \exp(u_w^T h)}$$

Where  $h$  is either:

- Average:  $h = \frac{1}{k+1} (v_d + \sum_{j=i-k}^{i-1} v_{w_j})$
- Concatenation:  $h = (v_d, \sum_{j=i-k}^{i-1} v_{w_j})U$

## Doc2Vec (PV-DBOW)

PV-DBOW probability:

$$P(w_i|d) = \frac{\exp(u_{w_i}^T v_d)}{\sum_{w \in V} \exp(u_w^T v_d)}$$

Where  $v_d$  is the document vector.

# Let's build Doc2Vec

[Jupyter notebook](#)

# Summary: Evolution of Word Embeddings

1. One-hot encoding → No semantics
2. TF-IDF → Document-level patterns
3. Word2Vec → Local context patterns
4. Doc2Vec → Document embeddings

# Assignment on GitHub: